

Machine Learning

Kernel Methods

Dariusz Hosseini

dariusz.hosseini@ucl.ac.uk
Department of Computer Science
University College London

Lecture Overview

- 1** Lecture Overview
- 2 The Kernel Trick
- 3 Mercer's Theorem
- 4 The Representer Theorem
- 5 Example: Ridge Regression
- 6 Summary
- 7 Appendix: Lagrange Duality

Lecture Overview

By the end of this lecture you should:

- 1 Know that **Kernel Methods** provide a mechanism for tractably accessing non-linear features for use with linear machine learning methods
- 2 Understand the relationship between **feature maps** and **kernels** via **Mercer's Theorem**
- 3 Understand when it is sensible to apply kernel methods via a consideration of the **Representer Theorem**

Lecture Overview

- 1 Lecture Overview
- 2 The Kernel Trick**
- 3 Mercer's Theorem
- 4 The Representer Theorem
- 5 Example: Ridge Regression
- 6 Summary
- 7 Appendix: Lagrange Duality

Overview

- Kernel methods seek to **implicitly** map our data into a **higher-dimensional feature space** so that **linear** machine learning algorithms can be used in this feature space
- Why?
 - Linear methods are well understood and often efficient
 - An implicit mapping allows us to access high, even infinite, dimensional feature spaces efficiently

Overview


- Usually kernel based methods will follow two steps:
 - 1 Map the data from the input space, \mathbb{X} , to a higher-dimensional **feature space**, \mathbb{V} , using some non-linear mapping, $\phi : \mathbb{X} \rightarrow \mathbb{V}$
 - 2 Run a linear machine learning algorithm in this new space
- This approach has wide applicability - many learning algorithms can be transformed into non-linear high dimensional methods, for example:
 - Ridge Regression
 - Support Vector Machine
 - PCA

Example

■ Input Data:

■ $\mathbf{x} = x_1 \in \mathbb{X}$

■ $\mathbb{X} \subseteq \mathbb{R}^1$




Example

■ Feature Mapped Data:

■ $\phi(\mathbf{x}) = [x_1, x_1^2]^T \in \mathbb{V}$

■ $\mathbb{V} \subseteq \mathbb{R}^2$




Example

■ Input Data:

■ $\mathbf{x} = [x_1, x_2]^T \in \mathbb{X}$

■ $\mathbb{X} \subseteq \mathbb{R}^2$




Example

■ Feature Mapped Data:

■ $\phi(\mathbf{x}) = [x_1^2, \sqrt{2}x_1x_2, x_2^2]^T \in \mathbb{V}$

■ $\mathbb{V} \subseteq \mathbb{R}^3$



Considerations: Why

- Are we justified in applying this strategy?
- Why should a linear relationship exist in a higher dimensional space?
- **Cover's Theorem** gives some intuition

Cover's Theorem

- Assume that our problem is that of binary classification
- The theorem states that the probability that n data points can be linearly separated in m dimensions is given by:

$$\begin{cases} 1, & n \leq m + 1 \\ \frac{1}{2^{n-1}} \sum_{i=0}^m \binom{n-1}{i}, & n \geq m + 1 \end{cases}$$

- Thus the probability of n instances being linearly separable increases with growing dimensionality m

Considerations: How

- How can we operate efficiently in a high dimensional space?
 - For example, consider **matrix inversion** in ridge regression for an infinite dimensional feature space
- We would like to be able to enjoy the benefits of feature mapping without having to perform that mapping or to perform manipulations in elevated dimensions **explicitly**
- The **Kernel Trick** let's us do exactly this

Kernel Trick

- Consider two input points, $\mathbf{x}^{(i)}, \mathbf{x}^{(j)} \in \mathbb{X}$
- Consider a feature mapping, $\phi : \mathbb{X} \rightarrow \mathbb{V}$
- We can define an **inner product** on \mathbb{V} , via the **kernel function**, κ , such that:

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \phi(\mathbf{x}^{(i)}) \cdot \phi(\mathbf{x}^{(j)})$$

Kernel Trick

- Let us consider the description of our learning algorithm for a moment:
- *'If an algorithm is described solely in terms of inner products, $\mathbf{x}^{(i)} \cdot \mathbf{x}^{(j)}$, in input space, then it can be lifted into feature space by replacing the occurrence of those inner products by $\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})$ ' [Rasmussen & Williams (2006)]*
- But what kind of algorithm looks like that?

Aside: Ridge Regression

- Recall the optimisation problem associated with Ridge Regression:

$$\min_{\mathbf{w}} \sum_{i=1}^n \left(y^{(i)} - \mathbf{w} \cdot \mathbf{x}^{(i)} \right)^2 + \lambda \|\mathbf{w}\|_2^2$$

- This is equivalent to:

$$\begin{aligned} \min_{\mathbf{w}} \quad & \sum_{i=1}^n \xi^{(i)2} + \lambda \|\mathbf{w}\|_2^2 \\ \text{subject to:} \quad & \xi^{(i)} = y^{(i)} - \mathbf{w} \cdot \mathbf{x}^{(i)} \end{aligned} \tag{1}$$

Aside: Ridge Regression & Lagrange Duality

- From the Appendix we see this problem can be solved by seeking the solution to its dual optimisation problem
- First we write the Lagrangian for problem (1):

$$\mathcal{L}(\xi, \mathbf{w}, \tilde{\alpha}) = \sum_{i=1}^n \xi^{(i)2} + \sum_{i=1}^n \tilde{\alpha}^{(i)} \left(y^{(i)} - \xi^{(i)} - \mathbf{w} \cdot \mathbf{x}^{(i)} \right) + \lambda \|\mathbf{w}\|_2^2$$

Here $\tilde{\alpha} = [\tilde{\alpha}^{(1)}, \dots, \tilde{\alpha}^{(n)}]$ are Lagrange Multipliers

- The dual objective can be written:

$$\mathcal{D}(\tilde{\alpha}) = \min_{\mathbf{w}, \xi} \mathcal{L}(\xi, \mathbf{w}, \tilde{\alpha})$$

Aside: Ridge Regression & Lagrange Duality

- This is an unconstrained optimisation which we can solve by seeking stationary points:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}} = - \sum_{i=1}^n \tilde{\alpha}^{(i)} \mathbf{x}^{(i)} + 2\lambda \mathbf{w}^* = 0 \quad \implies \quad \mathbf{w}^* = \frac{1}{2\lambda} \sum_{i=1}^n \tilde{\alpha}^{(i)} \mathbf{x}^{(i)}$$

$$\frac{\partial \mathcal{L}}{\partial \xi^{(i)}} = 2\xi^{(i)*} - \tilde{\alpha}^{(i)} = 0 \quad \implies \quad \xi^{(i)*} = \frac{\tilde{\alpha}^{(i)}}{2}$$

Aside: Ridge Regression & Lagrange Duality

- Substituting these expressions back into $\mathcal{D}(\tilde{\alpha})$ yields:

$$\begin{aligned}
 \mathcal{D}(\tilde{\alpha}) &= \frac{1}{4} \sum_{i=1}^n \tilde{\alpha}^{(i)2} + \sum_{i=1}^n \tilde{\alpha}^{(i)} y^{(i)} - \frac{1}{2} \sum_{i=1}^n \tilde{\alpha}^{(i)2} \\
 &\quad - \frac{1}{2\lambda} \sum_{i,j=1}^n \tilde{\alpha}^{(i)} \tilde{\alpha}^{(j)} \mathbf{x}^{(i)} \cdot \mathbf{x}^{(j)} + \lambda \left(\frac{1}{4\lambda^2} \left\| \sum_{i=1}^n \tilde{\alpha}^{(i)} \mathbf{x}^{(i)} \right\|_2^2 \right) \\
 &= -\frac{1}{4} \sum_{i=1}^n \tilde{\alpha}^{(i)2} + \sum_{i=1}^n \tilde{\alpha}^{(i)} y^{(i)} - \frac{1}{4\lambda} \sum_{i,j=1}^n \tilde{\alpha}^{(i)} \tilde{\alpha}^{(j)} \mathbf{x}^{(i)} \cdot \mathbf{x}^{(j)} \\
 &= -\lambda^2 \sum_{i=1}^n \alpha^{(i)2} + 2\lambda \sum_{i=1}^n \alpha^{(i)} y^{(i)} - \lambda \sum_{i,j=1}^n \alpha^{(i)} \alpha^{(j)} \mathbf{x}^{(i)} \cdot \mathbf{x}^{(j)}
 \end{aligned}$$

where: $\tilde{\alpha}^{(i)} = 2\lambda\alpha^{(i)}$

Aside: Ridge Regression & Dual Problem

- This leads to the following dual problem:

$$\begin{aligned} \max_{\alpha} \quad & -\lambda \|\alpha\|^2 + 2\alpha \cdot \mathbf{y} - \alpha^T (\mathbf{X}\mathbf{X}^T) \alpha \\ \text{or:} \quad & \max_{\alpha} \quad -\alpha^T (\mathbf{X}\mathbf{X}^T + \lambda \mathbf{I}) \alpha + 2\alpha \cdot \mathbf{y} \end{aligned}$$

- The solution to which is:

$$\begin{aligned} \mathbf{w}^* &= \mathbf{X}^T \alpha \\ \alpha &= (\mathbf{X}\mathbf{X}^T + \lambda \mathbf{I})^{-1} \mathbf{y} \end{aligned}$$

Kernel Trick: Revisited

- Again let us consider the description of our learning algorithm:
- *'If an algorithm is described solely in terms of inner products, $\mathbf{x}^{(i)} \cdot \mathbf{x}^{(j)}$, in input space, then it can be lifted into feature space by replacing the occurrence of those inner products by $\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})$ ' [Rasmussen & Williams (2006)]*
- For example, consider the dual problem of Ridge Regression:

$$\begin{aligned}
 & \max_{\alpha} \quad -\alpha^T(\mathbf{X}\mathbf{X}^T + \lambda\mathbf{I})\alpha + 2\alpha \cdot \mathbf{y} \\
 \implies & \max_{\{\alpha^{(i)}\}_{i=1}^n} \quad -\sum_{i,j=1}^n \alpha^{(i)} \alpha^{(j)} \mathbf{x}^i \cdot \mathbf{x}^j - \lambda \sum_{i=1}^n \alpha^{(i)2} + 2 \sum_{i=1}^n \alpha^{(i)} y^{(i)}
 \end{aligned}$$

Kernel Trick

- This is the **Kernel Trick**:

- It means that, for certain problems, we can enjoy the benefits of feature mapping without ever having to calculate $\phi(\mathbf{x})$ explicitly

- Instead we can calculate the mapping **implicitly** by calculating $\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \phi(\mathbf{x}^{(i)}) \cdot \phi(\mathbf{x}^{(j)})$

- In particular, it turns out that for many machine learning algorithms, when their associated optimisation problem is re-cast in its **dual form**, the inner product arises naturally:

- We shall investigate the characteristics of such problems later on

Lecture Overview

- 1 Lecture Overview
- 2 The Kernel Trick
- 3 Mercer's Theorem**
- 4 The Representer Theorem
- 5 Example: Ridge Regression
- 6 Summary
- 7 Appendix: Lagrange Duality

What is a Kernel?

- We now know that we can exploit the kernel trick to replace the inner products in (certain) learning algorithms with a kernel function
- And we would like to be able to work directly with such functions
- But how do we know if a particular function, κ , really has a feature map, ϕ , associated with it such that,
$$\kappa(\mathbf{x}^{(i)} \cdot \mathbf{x}^{(j)}) = \phi(\mathbf{x}^{(i)}) \cdot \phi(\mathbf{x}^{(j)})?$$
- And what characterises a valid function, κ ?

Mercer's Theorem

■ Definition:

The **Gram Matrix**, \mathbf{K} , is an $n \times n$ matrix with elements, $\{\mathbf{K}_{ij} = \kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})\}_{i=1, j=1}^{n, n}$ where $\mathbf{x}^{(i)} \in \mathbb{R}^m$

■ Definition:

$\kappa : \mathbb{R}^m \times \mathbb{R}^m \rightarrow \mathbb{R}$ is positive semidefinite (psd) if it is symmetric and if the Gram Matrix, \mathbf{K} , is psd for all n and for all $\mathbf{x}^{(i)} \in \mathbb{R}^m$

Mercer's Theorem

■ Mercer's Theorem:

A kernel function, κ , is psd if and only if:

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \phi(\mathbf{x}^{(i)}) \cdot \phi(\mathbf{x}^{(j)}) \quad \text{where} \quad \mathbf{x}^{(i)}, \mathbf{x}^{(j)} \in \mathbb{R}^m$$

For some feature map, $\phi : \mathbb{R}^m \rightarrow \mathbb{W}$ (and Hilbert Space \mathbb{W})

■ We call a kernel which satisfies this theorem a **Mercer Kernel**

Mercer's Theorem: Aside

■ Let us prove the only if:

■ If:

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \phi(\mathbf{x}^{(i)}) \cdot \phi(\mathbf{x}^{(j)})$$

■ Then:

$$\begin{aligned}\mathbf{c}^T \mathbf{K} \mathbf{c} &= \sum_{i,j=1}^n c_i c_j \kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \left(\sum_{i=1}^n c_i \phi(\mathbf{x}^{(i)}) \right) \cdot \left(\sum_{j=1}^n c_j \phi(\mathbf{x}^{(j)}) \right) \\ &= \|\mathbf{c}_i \phi(\mathbf{x}^{(i)})\|_2^2 \geq 0\end{aligned}$$

This holds for all $\mathbf{c} \in \mathbb{R}^n$ where c_i is the i -th element of \mathbf{c}

Polynomial Kernel

- The **polynomial kernel** can be used to compute a compact version of the polynomial feature mapping up to degree p :

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = (1 + \mathbf{x}^{(i)} \cdot \mathbf{x}^{(j)})^p$$

- Let's try expanding for $p = 2$ and $\mathbf{x} \in \mathbb{R}^2$:

$$\begin{aligned}\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) &= (1 + x_1^{(i)} x_1^{(j)} + x_2^{(i)} x_2^{(j)})^2 \\ &= 1 + 2x_1^{(i)} x_1^{(j)} + 2x_2^{(i)} x_2^{(j)} + (x_1^{(i)} x_1^{(j)})^2 + (x_2^{(i)} x_2^{(j)})^2 + 2x_1^{(i)} x_2^{(i)} x_1^{(j)} x_2^{(j)}\end{aligned}$$

- In other words our use of the kernel is equivalent to a mapping into 6-dimensional feature space:

$$\phi(\mathbf{x}) = [1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, x_2^2, \sqrt{2}x_1 x_2]^T$$

Radial Basis Function Kernel

- The **RBF kernel** is a flexible measure of similarity parameterised by a **bandwidth** hyperparameter, σ^2 :

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \exp\left(-\frac{\|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|_2^2}{2\sigma^2}\right)$$

- The RBF kernel actually gives access to an infinite dimensional feature space.

Document Kernel

- Kernels can be constructed for a variety of data types (graphs, categorical data, etc.)
- The **cosine similarity kernel** can be used as a **document kernel** for the **bag-of-words** representation:

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \frac{\mathbf{x}^{(i)} \cdot \mathbf{x}^{(j)}}{\|\mathbf{x}^{(i)}\|_2 \|\mathbf{x}^{(j)}\|_2}$$

- Here the input attributes are all counts, and so the measure falls in the range $[0, 1]$

Recap

- We now know that:
 - **Mercer's Theorem** allows us to specify a **feature map** implicitly if we specify a **Mercer kernel**
 - And the **Kernel Trick** makes clear that if our learning problem can be written in terms of dot products of the input vectors then we can move to this kernel-defined feature space by simply replacing the dot products with the corresponding kernel mapped output
 - ...But how do we know if a particular learning problem can be written in this way...
 - ...For this we need to consider the **Representer Theorem**

Lecture Overview

- 1 Lecture Overview
- 2 The Kernel Trick
- 3 Mercer's Theorem
- 4 The Representer Theorem**
- 5 Example: Ridge Regression
- 6 Summary
- 7 Appendix: Lagrange Duality

Notation

- **Training Sample:**

$$\{(\mathbf{x}^{(i)}, y^{(i)}) | \mathbf{x}^{(i)} \in \mathbb{X} \subseteq \mathbb{R}^m, y^{(i)} \in \mathbb{R}\}_{i=1}^n$$

- **Feature Map:**

$$\phi : \mathbb{R}^m \rightarrow \mathbb{R}^k$$

- **Weight Vector:**

$$\mathbf{w} \in \mathbb{R}^k$$

- **Linear Prediction Function:**

$$f(\mathbf{x}) = \mathbf{w} \cdot \phi(\mathbf{x})$$

- **Loss Function:**

$$\tilde{L} : (y, f(\mathbf{x})) \mapsto \tilde{L}(y, f(\mathbf{x})) \in \mathbb{R}$$

- **And some function:**

$$\Omega : \mathbf{w} \mapsto \Omega(\mathbf{w}) \in \mathbb{R}$$

Representer Theorem: Sketch

- For a regularised loss function, L , defined such that:

$$L(\mathbf{w}) = \sum_{i=1}^n \tilde{L}(y^{(i)}, \mathbf{w} \cdot \phi(\mathbf{x}^{(i)})) + \Omega(\mathbf{w})$$

- If and only if $\Omega(\mathbf{w})$ is a non-decreasing function of $\|\mathbf{w}\|_2^2$ then, if \mathbf{w}^* minimises L , it admits the following representation:

$$\mathbf{w}^* = \sum_{i=1}^n \alpha_i \phi(\mathbf{x}^{(i)}) \quad \text{where: } \alpha_i \in \mathbb{R}$$

- And therefore for a novel test point, \mathbf{z} :

$$f(\mathbf{z}) = \mathbf{w} \cdot \phi(\mathbf{z}) = \sum_{i=1}^n \alpha_i \kappa(\mathbf{x}^{(i)}, \mathbf{z})$$

Representer Theorem

- This is a powerful result
 - For learning algorithms whose optimisation problem satisfies the conditions of the representer theorem...
 - ...We need never directly seek \mathbf{w}^* (whose dimensions are possibly infinite)
 - Instead we need only seek the n parameters which characterise α
 - Furthermore, when evaluating a test point we need never explicitly calculate $\phi(\mathbf{z})$
 - Instead we need only sum over a weighted set of n kernel function outputs, evaluated as the contraction of \mathbf{z} and each of the training points $\{\mathbf{x}^{(i)}\}_{i=1}^n$

Representer Theorem: Sparsity

- Note that a **sparsity-inducing** regulariser such as $\Omega(\mathbf{w}) = \|\mathbf{w}\|_1$, as in the **LASSO**, does not satisfy the conditions of the Representer Theorem
- This explains why we cannot use the kernel trick in conjunction with this sort of **feature selection**
- Sparsity comes at a price!

Lecture Overview

- 1 Lecture Overview
- 2 The Kernel Trick
- 3 Mercer's Theorem
- 4 The Representer Theorem
- 5 Example: Ridge Regression**
- 6 Summary
- 7 Appendix: Lagrange Duality

Ridge Regression: Primal Form

- Recall our **Ridge Regression** optimisation problem, where we have elevated our input attributes to a k -dimensional feature space via the map, $\phi : \mathbb{R}^m \rightarrow \mathbb{R}^k$:

$$\underset{\mathbf{w}}{\operatorname{argmin}} \quad \sum_{i=1}^n \left(y^{(i)} - \mathbf{w} \cdot \phi(\mathbf{x}^{(i)}) \right)^2 + \lambda \|\mathbf{w}\|_2^2$$

- Here the so-called **primal solution** is generated by the **revised normal equations**:

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y} \quad (2)$$

Where \mathbf{X} is the $n \times k$ feature mapped **design matrix**,
 $\mathbf{X} = [\phi(\mathbf{x}^{(1)}), \phi(\mathbf{x}^{(2)}), \dots, \phi(\mathbf{x}^{(n)})]^T$

- And the associated prediction for a novel test point \mathbf{z} is:

$$f(\mathbf{z}) = \mathbf{w} \cdot \phi(\mathbf{z}) = \mathbf{y}^T \mathbf{X} (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \phi(\mathbf{z})$$

Ridge Regression: Dual Form

- Left multiply equation (1) by $(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})$:

$$\begin{aligned}\mathbf{X}^T \mathbf{X} \mathbf{w} + \lambda \mathbf{w} &= \mathbf{X}^T \mathbf{y} \\ \mathbf{w} &= \lambda^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{X} \mathbf{w}) \\ &= \mathbf{X}^T \boldsymbol{\alpha}\end{aligned}\tag{3}$$

- Here:

$$\begin{aligned}\boldsymbol{\alpha} &= \lambda^{-1} (\mathbf{y} - \mathbf{X} \mathbf{w}) \\ \lambda \boldsymbol{\alpha} &= (\mathbf{y} - \mathbf{X} \mathbf{w}) \\ \lambda \boldsymbol{\alpha} &= (\mathbf{y} - \mathbf{X} \mathbf{X}^T \boldsymbol{\alpha}) \quad \text{From equation (3)} \\ \boldsymbol{\alpha} &= (\mathbf{X} \mathbf{X}^T + \lambda \mathbf{I})^{-1} \mathbf{y}\end{aligned}$$

This is the so-called **dual solution**

Ridge Regression: Dual Form

- And the associated prediction for a novel test point \mathbf{z} is:

$$\begin{aligned}f(\mathbf{z}) &= \mathbf{w} \cdot \phi(\mathbf{z}) \\&= \mathbf{y}^T (\mathbf{X}\mathbf{X}^T + \lambda \mathbf{I})^{-1} \mathbf{X} \phi(\mathbf{z}) \\&= \sum_{i=1}^n \alpha_i \phi(\mathbf{x}^{(i)}) \cdot \phi(\mathbf{z})\end{aligned}$$

- In fact the same result follows from:

- The **Representer Theorem**
- The **Lagrangian Dual solution**...as we observed earlier

Ridge Regression: Kernel Version

- We may re-write this using kernels as:

$$f(\mathbf{z}) = \sum_{i=1}^n \alpha_i \kappa(\mathbf{x}^{(i)}, \mathbf{z})$$

Where:

$$\begin{aligned}\boldsymbol{\alpha} &= (\mathbf{X}\mathbf{X}^T + \lambda\mathbf{I})^{-1}\mathbf{y} \\ (\mathbf{X}\mathbf{X}^T)_{ij} &= \kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)})\end{aligned}$$

- We see that $\phi(\mathbf{x})$ is never explicitly evaluated either within the calculation of α_i or within the kernelised prediction function

Ridge Regression: The Value of the Kernel Trick

- We note that the primal solution involves an inversion of the **design matrix**, $\mathbf{X}^T \mathbf{X}$, an operation of complexity $\mathcal{O}(k^3)$
...and prediction involves an operation of complexity $\mathcal{O}(k)$
- Meanwhile the dual solution involves an inversion of the **Gram matrix**, $\mathbf{X}\mathbf{X}^T$, an operation of complexity $\mathcal{O}(n^3)$
...and prediction involves an operation of complexity $\mathcal{O}(n|\kappa|)$
(where $|\kappa|$ is the number of operations involved in the calculation of an inner product)
- Thus, for $k \gg n$ the value of the dual form and the kernel trick becomes apparent

Lecture Overview

- 1 Lecture Overview
- 2 The Kernel Trick
- 3 Mercer's Theorem
- 4 The Representer Theorem
- 5 Example: Ridge Regression
- 6 Summary**
- 7 Appendix: Lagrange Duality

Summary

- 1 We can attempt to enhance the **linear** algorithms of machine learning by affecting a **feature map** of the input attributes in order to elevate them to a higher dimensional **non-linear** space
- 2 This is a sensible and a tractable strategy if we can employ the **Kernel Trick** and work only via the dot product of feature mapped data points
- 3 The kernel trick is manifest iff a **Representer Theorem** holds for our optimisation problem
- 4 In this case, **Mercer's Theorem** offers us a mechanism for defining a valid feature space via **Mercer Kernels**, without the need to ever explicitly define the feature mapping

Lecture Overview

- 1 Lecture Overview
- 2 The Kernel Trick
- 3 Mercer's Theorem
- 4 The Representer Theorem
- 5 Example: Ridge Regression
- 6 Summary
- 7 Appendix: Lagrange Duality**

Multiple Constraints: Problem



$$\begin{array}{ll} \min_{\mathbf{x} \in \mathbb{R}^n} & f(\mathbf{x}) \\ \text{subject to:} & \left\{ \begin{array}{l} \{g^{(i)}(\mathbf{x}) \leq 0\}_{i=1}^m \\ \{h^{(j)}(\mathbf{x}) = 0\}_{j=1}^p \end{array} \right. \end{array}$$

Multiple Constraints: Lagrangian

- We express the Lagrangian as:

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}, \boldsymbol{\mu}) = f(\mathbf{x}) + \sum_{i=1}^m \mu^{(i)} g^{(i)}(\mathbf{x}) + \sum_{j=1}^p \lambda^{(j)} h^{(j)}(\mathbf{x})$$

Where:

$$\boldsymbol{\lambda} = [\lambda^{(1)}, \dots, \lambda^{(p)}]^T, \{\lambda^{(j)} \in \mathbb{R}\}_{j=1}^p;$$

$$\boldsymbol{\mu} = [\mu^{(1)}, \dots, \mu^{(m)}]^T, \{\mu^{(i)} \in \mathbb{R}_{\geq 0}\}_{i=1}^m;$$

are Lagrange multipliers

Multiple Constraints: Problem Reformulation

- And we can solve our problem by seeking stationary solutions $(\mathbf{x}^*, \{\mu^{(i)*}\}, \{\lambda^{(j)*}\})$ which satisfy the following:

$$\begin{aligned} & \nabla_{\mathbf{x}} \mathcal{L} = \mathbf{0} \\ \text{subject to: } & \left\{ \begin{array}{l} \{g^{(i)}(\mathbf{x}) \leq 0\}_{i=1}^m, \{h^{(j)}(\mathbf{x}) = 0\}_{j=1}^p \\ \{\mu^{(i)} \geq 0\}_{i=1}^m \\ \{\mu^{(i)} g^{(i)}(\mathbf{x}) = 0\}_{i=1}^m \end{array} \right. \end{aligned}$$

Duality: Primal Problem

- The original problem is sometimes known as the **primal problem**, and its variables, \mathbf{x} , are known as the **primal variables**
- It is equivalent to the following formulation:

$$\min_{\mathbf{x}} \left[\max_{\lambda, \mu \geq 0} \mathcal{L}(\mathbf{x}, \lambda, \mu) \right]$$

- Here the bracketed term is known as the **primal objective** function

Duality: Barrier Function

- We can re-write the primal objective as follows:

$$\max_{\lambda, \mu \geq 0} \mathcal{L}(\mathbf{x}, \lambda, \mu) = f(\mathbf{x}) + \max_{\lambda, \mu \geq 0} \left[\sum_{i=1}^m \mu^{(i)} g^{(i)}(\mathbf{x}) + \sum_{j=1}^p \lambda^{(j)} h^{(j)}(\mathbf{x}) \right]$$

- Here the second term gives rise to a **barrier function** which enforces the constraints as follows:

$$\max_{\lambda, \mu \geq 0} \left[\sum_{i=1}^m \mu^{(i)} g^{(i)}(\mathbf{x}) + \sum_{j=1}^p \lambda^{(j)} h^{(j)}(\mathbf{x}) \right] = \begin{cases} 0 & \text{if } \mathbf{x} \text{ is feasible} \\ \infty & \text{if } \mathbf{x} \text{ is infeasible} \end{cases}$$

Duality: Minimax Inequality

- In order to make use of this barrier function formulation, we will need the **minimax inequality**:

$$\max_{\mathbf{y}} \min_{\mathbf{x}} \phi(\mathbf{x}, \mathbf{y}) \leq \min_{\mathbf{x}} \max_{\mathbf{y}} \phi(\mathbf{x}, \mathbf{y})$$

- **Proof:**

$$\min_{\mathbf{x}} \phi(\mathbf{x}, \mathbf{y}) \leq \phi(\mathbf{x}, \mathbf{y}) \quad \forall \mathbf{x}, \mathbf{y}$$

This is true for all \mathbf{y} , therefore, in particular the following is true:

$$\max_{\mathbf{y}} \min_{\mathbf{x}} \phi(\mathbf{x}, \mathbf{y}) \leq \max_{\mathbf{y}} \phi(\mathbf{x}, \mathbf{y}) \quad \forall \mathbf{x}$$

This is true for all \mathbf{x} , therefore, in particular the following is true:

$$\max_{\mathbf{y}} \min_{\mathbf{x}} \phi(\mathbf{x}, \mathbf{y}) \leq \min_{\mathbf{x}} \max_{\mathbf{y}} \phi(\mathbf{x}, \mathbf{y})$$

Duality: Weak Duality

- We can now introduce the concept of **weak duality**:

$$\min_{\mathbf{x}} \left[\max_{\lambda, \mu \geq 0} \mathcal{L}(\mathbf{x}, \lambda, \mu) \right] \geq \max_{\lambda, \mu \geq 0} \left[\min_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \lambda, \mu) \right]$$

- Here the bracketed term on the right hand side is known as the **dual objective** function, $\mathcal{D}(\lambda, \mu)$
- If we can solve the right hand side of the inequality then we have a lower bound on the solution of our optimisation problem

Duality: Weak Duality

- And often the RHS side of the inequality is an **easier** problem to solve, because:
 - $\min_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \lambda, \mu)$ is an **unconstrained** optimisation problem for a given value of (λ, μ) ...
 - ...And if solving this problem is not hard then the overall problem is not hard to solve because:
 - $\max_{\lambda, \mu \geq 0} [\min_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \lambda, \mu)]$ is a maximisation problem over a set of affine functions - thus it is a **concave maximisation** problem or equivalently a **convex minimisation** problem, and we know that such problems can be efficiently solved
 - Note that this is true regardless of whether $f, g^{(i)}, h^{(j)}$ are nonconvex

Duality: Strong Duality

- For certain classes of problems which satisfy **constraint qualifications** we can go further and **strong duality** holds:

$$\min_{\mathbf{x}} \left[\max_{\lambda, \mu \geq 0} \mathcal{L}(\mathbf{x}, \lambda, \mu) \right] = \max_{\lambda, \mu \geq 0} \left[\min_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \lambda, \mu) \right]$$

- There are several different constraint qualifications. One is **Slater's Condition** which holds for **convex optimisation** problems
- Recall, these are problems for which f is convex and $g^{(i)}, h^{(j)}$ are convex sets
- For problems of this type we may seek to solve the **dual optimisation** problem:

$$\max_{\lambda, \mu \geq 0} \left[\min_{\mathbf{x}} \mathcal{L}(\mathbf{x}, \lambda, \mu) \right]$$

Duality: Strong Duality

- Another reason for adopting the dual optimisation approach to solving constrained optimisation problems is based on dimensionality:
- If the dimensionality of the dual variables, $(m + p)$, is less than the dimensionality of the primal variables, n , then dual optimisation often offers a more efficient route to solutions
- This is of particular importance if we are dealing with infinite dimensional primal variables