

Machine Learning

The Non-Linear SVM

Dariusz Hosseini

dariusz.hosseini@ucl.ac.uk
Department of Computer Science
University College London

Lecture Overview

- 1** Lecture Overview
- 2 Support Vector Classification
- 3 Multi-Class Support Vector Classification
- 4 Support Vector Regression
- 5 Summary

Lecture Overview

By the end of this lecture you should:

- 1 Understand how we can extend the **Linear SVM** using **Kernels** to create the **Non-Linear SVM**
- 2 Know how we can extend the **binary** support vector classifier to enable us to tackle **multi-class** learning problems
- 3 Know that SVMs can be used for **regression** as well as classification

Lecture Overview

- 1 Lecture Overview
- 2 Support Vector Classification**
- 3 Multi-Class Support Vector Classification
- 4 Support Vector Regression
- 5 Summary

Linear SVM: Optimisation Problem

- Recall the original linear SVM problem:

$$\min_{\mathbf{w}, b} \quad \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \max(0, y^{(i)} (\mathbf{w} \cdot \mathbf{x}^{(i)} + b)) \quad (1)$$

- And the associated **dual problem**:

$$\max_{\alpha} \quad \sum_{i=1}^n \alpha^{(i)} - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha^{(i)} \alpha^{(j)} y^{(i)} y^{(j)} \mathbf{x}^{(i)} \cdot \mathbf{x}^{(j)} \quad (2)$$

$$\begin{aligned} \text{subject to:} \quad & \sum_{i=1}^n \alpha^{(i)} y^{(i)} = 0 \\ & 0 \leq \alpha^{(i)} \leq C \end{aligned}$$

Linear SVM: Optimisation Solution & Prediction

- Recall that we could express our solution to the problem as:

$$\mathbf{w}^* = \sum_{i \in \mathcal{SV}} \alpha^{(i)*} y^{(i)} \mathbf{x}^{(i)}$$
$$b^* = \frac{1}{|\widetilde{\mathcal{SV}}|} \sum_{i \in \widetilde{\mathcal{SV}}} \left(y^{(i)} - \sum_{j \in \widetilde{\mathcal{SV}}} \alpha^{(j)*} y^{(j)} \mathbf{x}^{(j)} \cdot \mathbf{x}^{(i)} \right)$$


Where \mathcal{SV} is the set of support vectors, and $\widetilde{\mathcal{SV}}$ is the set of support vectors for which $0 < \alpha^{(i)} < C$

- And predict the class of a novel test point, \mathbf{z} , as:

$$f(\mathbf{z}) = \text{sgn}(\mathbf{w}^* \cdot \mathbf{z} + b^*)$$
$$= \text{sgn} \left(\sum_{i \in \mathcal{SV}} \alpha^{(i)*} y^{(i)} \mathbf{x}^{(i)} \cdot \mathbf{z} + b^* \right)$$

The Failure Case

- Recall that this works well for both **separable** and **noisy** settings...if the boundary is **linear**
- But not well for **non-linear boundaries**:



- Can we use **kernel methods** to enhance the algorithm?

Kernel Methods: Recap

■ Kernel Trick:

- If the dependency of input attributes within our algorithm can be expressed solely in terms of **inner products** between the input vectors...
- ...Then we can replace all such inner products with an appropriate **kernel function** output

■ Mercer's Theorem:

- Gives us criteria for the **validity** of kernels which we may use in the kernel trick such that they will implicitly express a valid feature mapping

■ Representer Theorem:

- Gives necessary and sufficient conditions that the form of an optimisation problem must take such that it will admit the kernel trick

Representer Theorem: Recap

- For a regularised loss function, L , defined such that:

$$L(\mathbf{w}) = \sum_{i=1}^n \tilde{L}(y^{(i)}, \mathbf{w} \cdot \phi(\mathbf{x}^{(i)})) + \Omega(\mathbf{w})$$

- If and only if $\Omega(\mathbf{w})$ is a non-decreasing function of $\|\mathbf{w}\|_2$ then, if \mathbf{w}^* minimises L , it admits the following representation:

$$\mathbf{w}^* = \sum_{i=1}^n \tilde{\alpha}_i \phi(\mathbf{x}^{(i)}) \quad \text{where: } \tilde{\alpha}_i \in \mathbb{R}$$

- And therefore for a novel test point, \mathbf{z} :

$$f(\mathbf{z}) = \mathbf{w} \cdot \phi(\mathbf{z}) = \sum_{i=1}^n \tilde{\alpha}_i \kappa(\mathbf{x}^{(i)}, \mathbf{z})$$

Representer Theorem: SVM

- We note that the **SVM learning problem** given by expression (1), satisfies the requirement of the **Representer Theorem**
- Thus we can state that:

$$\mathbf{w}^* = \sum_{i=1}^n \tilde{\alpha}_i \mathbf{x}^{(i)}$$

- Substituting this back into the original learning problem will lead to a dual form, from which we see that:

$$\tilde{\alpha}_i = \alpha^{(i)} y^{(i)}$$

- And so we are led back to the **dual solution** of expression (2) by a different route

Representer Theorem: SVM

- So we can apply the kernel trick to the SVM:

$$\begin{aligned}\phi(\mathbf{x}) &\longleftarrow \mathbf{x} \\ \kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) &\longleftarrow \mathbf{x}^{(i)} \cdot \mathbf{x}^{(j)}\end{aligned}$$

Here:

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \phi(\mathbf{x}^{(i)}) \cdot \phi(\mathbf{x}^{(j)})$$

Non-Linear SVM: Optimisation Problem

- The Non-Linear SVM problem becomes:

$$\min_{\mathbf{w}, b} \quad \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \max(0, y^{(i)} (\mathbf{w} \cdot \phi(\mathbf{x}^{(i)}) + b)) \quad (3)$$

- With the associated **dual problem**:

$$\max_{\alpha} \quad \sum_{i=1}^n \alpha^{(i)} - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n \alpha^{(i)} \alpha^{(j)} y^{(i)} y^{(j)} \kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) \quad (4)$$

$$\begin{aligned} \text{subject to:} \quad & \sum_{i=1}^n \alpha^{(i)} y^{(i)} = 0 \\ & 0 \leq \alpha^{(i)} \leq C \end{aligned}$$

Non-Linear SVM: Optimisation Solution & Prediction

- And the solution to the non-linear problem is:

$$\mathbf{w}^* = \sum_{i \in \mathcal{SV}} \alpha^{(i)*} y^{(i)} \phi(\mathbf{x}^{(i)})$$

$$b^* = \frac{1}{|\widetilde{\mathcal{SV}}|} \sum_{i \in \widetilde{\mathcal{SV}}} \left(y^{(i)} - \sum_{j \in \widetilde{\mathcal{SV}}} \alpha^{(j)*} y^{(j)} \kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) \right)$$

- And we can predict the class of a novel test point, \mathbf{z} , as:

$$\begin{aligned} f(\mathbf{z}) &= \text{sgn}(\mathbf{w}^* \cdot \phi(\mathbf{z}) + b^*) \\ &= \text{sgn} \left(\sum_{i \in \mathcal{SV}} \alpha^{(i)*} y^{(i)} \kappa(\mathbf{x}^{(i)}, \mathbf{z}) + b^* \right) \end{aligned}$$


Failure Case: RBF Solution

- Let's return to the failure case, but now we'll attempt to learn a **boundary** in the **feature space** defined by the **RBF kernel**:

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \exp\left(-\gamma\|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|^2\right)$$

- Recall that the RBF kernel is associated with an ∞ -order polynomial feature map, so it has the capacity to learn **complex boundaries**

Failure Case: RBF Solution




- Soft-margin classifier with RBF Kernel ($\gamma = 0.02$)

Hyperparameters

- How do we set the **hyperparameters** C and γ ?
- **Cross-validation** is usually employed
- Increasing C leads to:
 - Less tolerance of errors
 - More complex boundaries
- Increasing γ leads to:
 - Sharply peaked similarity measure
 - So each support vector becomes only locally influential
 - And we obtain more complex boundaries

Hyperparameters



Points to Remember

■ **Linearity**

- SVMs are a linear, maximal margin technique

■ **Kernel Trick**

- The use of kernels makes SVMs extremely flexible

■ **Sparsity**

- Once trained we need only retain the support vectors

■ **Convexity**

- The hinge loss and margin maximisation make the optimisation convex

Lecture Overview

- 1 Lecture Overview
- 2 Support Vector Classification
- 3 Multi-Class Support Vector Classification**
- 4 Support Vector Regression
- 5 Summary


SVMs with Multiple Classes?

- The original SVM is designed for **binary** classification
- We can extend the idea of a separating hyperplane to the case where we are attempting to classify **multiple classes**
- Two of the most popular approaches are:
 - **One-Versus-One** (OvO)
 - **One-Versus-All** (OvA)


One-versus-One

- The OvO approach for $K > 2$ classes constructs $\frac{K(K-1)}{2}$ SVMs, each of which compares a pair of classes
- A new instance is then classified using each classifier
- A tally is kept and the final classification is obtained by **majority vote**


One-versus-One



One-versus-One




One-versus-One



\times	f_1	f_2	f_3
●	✓	✓	
■			✓
▲			

One-versus-One



\times	f_1	f_2	f_3
●	✓	✓	
■			
▲			✓

$+$	f_1	f_2	f_3
●	✓		
■		✓	
▲			✓

One-versus-One

■ Pros:

- Easy to train classifiers (provided they are equally weighted)


■ Cons:

- Many classifiers to train
- Regions of ambiguity


One-versus-All

- The OvA approach constructs K SVMs, each of which compares the instances of one class to instances of *all* other classes
- A new instance is then classified by calculating the discriminant function for each of the K classifiers
- Final classification is obtained by assigning according to a positive classification, or, in the event of ties, according to the **maximal discriminant**


One-versus-All



One-versus-All




One-versus-All



\times	f_1	f_2	f_3
●		-1	
▲			-1
■	+1		

One-versus-All



\times	f_1	f_2	f_3
●		-1	
▲			-1
■	+1		

$+$	f_1	f_2	f_3
●		-1	
▲			-1
■	-1		

One-versus-All

■ Pros:

- Few classifiers to train

■ Cons:

- Class imbalance
- Regions of ambiguity
- Scaling of discriminant function needs to be tuned (distance to different hyperplanes is not measured on the same scale!)

A Consistent Approach

- The regions of ambiguity occur because OvO and OvA are just **heuristics**
- Each binary classifier does not know that we use its output prediction for a multi-class prediction - this might lead to sub-optimal results
- It is better to specify the complete task initially and seek to tackle the problem whole
- How can we do this?

K Class Discriminant

- We can, for example, specify a K class discriminant, which consists of K linear functions of the form:

$$f_i(\mathbf{x}) = \mathbf{w}_i \cdot \mathbf{x} + w_{i0}$$

Here i is a class index running from 1 to K

- Then we assign a point to a class associated with k if $f_k(\mathbf{x}) > f_j(\mathbf{x})$ for all $j \neq k$
- This results in **decision boundaries** between k and j given by:

$$\begin{aligned} f_k(\mathbf{x}) &= f_j(\mathbf{x}) \\ \implies (\mathbf{w}_k - \mathbf{w}_j) \cdot \mathbf{x} + (w_{k0} - w_{j0}) &= 0 \end{aligned}$$

- Will this still give rise to regions of ambiguity?


Non-Ambiguity of the K Class Discriminant

- No, because each of the decision regions defined by the discriminant functions is **convex**:
- Consider two points \mathbf{x}_A and \mathbf{x}_B which both lie in the decision region associated with k , R_k
- Then any point, $\tilde{\mathbf{x}}$, that lies on the line connecting \mathbf{x}_A and \mathbf{x}_B can be expressed as:

$$\tilde{\mathbf{x}} = \lambda \mathbf{x}_A + (1 - \lambda) \mathbf{x}_B$$

Here $0 \leq \lambda \leq 1$

Non-Ambiguity of the K Class Discriminant



Non-Ambiguity of the K Class Discriminant

- Using our decision functions we can write:

$$\begin{aligned}f_k(\tilde{\mathbf{x}}) &= \mathbf{w}_k \cdot \tilde{\mathbf{x}} + w_{k0} \\&= \mathbf{w}_k \cdot (\lambda \mathbf{x}_A + (1 - \lambda) \mathbf{x}_B) + (\lambda w_{k0} + (1 - \lambda) w_{k0}) \\&= \lambda f_k(\mathbf{x}_A) + (1 - \lambda) f_k(\mathbf{x}_B)\end{aligned}$$

- Because \mathbf{x}_A and \mathbf{x}_B lie in R_k , then $f_k(\mathbf{x}_A) > f_j(\mathbf{x}_A)$ and $f_k(\mathbf{x}_B) > f_j(\mathbf{x}_B)$ for all $j \neq k$
- Therefore $f_k(\tilde{\mathbf{x}}) > f_j(\tilde{\mathbf{x}})$ for all $j \neq k$, and so $\tilde{\mathbf{x}}$ also lies in R_k , and R_k is convex
- And this rules out any ambiguous regions, which would necessarily be non-convex

Aside: Softmax Regression

- How do we motivate such a K class discriminant?
- **Softmax Regression** provides one way
- This is a **probabilistic approach**, which is the multinomial generalisation of **logistic regression**
- We begin by noting that for **multinomial classification** and **misclassification loss** we can, w.l.o.g, prove the following **Bayes Optimal classifier**, f^* :

$$f^*(\mathbf{x}) = k \quad \text{if:} \quad p_y(y = k|\mathbf{x}) > p_y(y = j|\mathbf{x}) \quad \forall j \neq k$$

Aside: Softmax Regression

- Now, we make a set of modelling assumptions: We take one of the classes, say K , as the reference class, and then make the following assumption:

Aside: Softmax Regression

- Now, we make a set of modelling assumptions: We take one of the classes, say K , as the reference class, and then make the following assumption:

$$\log \left(\frac{p_y(y = i|\mathbf{x})}{p_y(y = K|\mathbf{x})} \right) = (\mathbf{w}_i - \mathbf{w}_K) \cdot \mathbf{x} + (w_{i0} - w_{K0})$$

Aside: Softmax Regression

- Now, we make a set of modelling assumptions: We take one of the classes, say K , as the reference class, and then make the following assumption:

$$\begin{aligned}\log \left(\frac{p_y(y = i|\mathbf{x})}{p_y(y = K|\mathbf{x})} \right) &= (\mathbf{w}_i - \mathbf{w}_K) \cdot \mathbf{x} + (w_{i0} - w_{K0}) \\ \left(\frac{p_y(y = i|\mathbf{x})}{p_y(y = K|\mathbf{x})} \right) &= \exp((\mathbf{w}_i - \mathbf{w}_K) \cdot \mathbf{x} + (w_{i0} - w_{K0}))\end{aligned}\quad (5)$$

Aside: Softmax Regression

- Now, we make a set of modelling assumptions: We take one of the classes, say K , as the reference class, and then make the following assumption:

$$\log \left(\frac{p_y(y = i|\mathbf{x})}{p_y(y = K|\mathbf{x})} \right) = (\mathbf{w}_i - \mathbf{w}_K) \cdot \mathbf{x} + (w_{i0} - w_{K0})$$
$$\left(\frac{p_y(y = i|\mathbf{x})}{p_y(y = K|\mathbf{x})} \right) = \exp((\mathbf{w}_i - \mathbf{w}_K) \cdot \mathbf{x} + (w_{i0} - w_{K0})) \quad (5)$$

$$\sum_{i=1}^{K-1} \left(\frac{p_y(y = i|\mathbf{x})}{p_y(y = K|\mathbf{x})} \right) = \frac{1 - p_y(y = K|\mathbf{x})}{p_y(y = K|\mathbf{x})}$$

Aside: Softmax Regression

- Now, we make a set of modelling assumptions: We take one of the classes, say K , as the reference class, and then make the following assumption:

$$\log \left(\frac{p_y(y = i|\mathbf{x})}{p_y(y = K|\mathbf{x})} \right) = (\mathbf{w}_i - \mathbf{w}_K) \cdot \mathbf{x} + (w_{i0} - w_{K0})$$
$$\left(\frac{p_y(y = i|\mathbf{x})}{p_y(y = K|\mathbf{x})} \right) = \exp((\mathbf{w}_i - \mathbf{w}_K) \cdot \mathbf{x} + (w_{i0} - w_{K0})) \quad (5)$$

$$\sum_{i=1}^{K-1} \left(\frac{p_y(y = i|\mathbf{x})}{p_y(y = K|\mathbf{x})} \right) = \frac{1 - p_y(y = K|\mathbf{x})}{p_y(y = K|\mathbf{x})}$$
$$= \sum_{i=1}^{K-1} \exp((\mathbf{w}_i - \mathbf{w}_K) \cdot \mathbf{x} + (w_{i0} - w_{K0}))$$

Aside: Softmax Regression

■ So:

$$p_y(y = K|\mathbf{x}) = \frac{1}{1 + \sum_{i=1}^{K-1} \exp((\mathbf{w}_i - \mathbf{w}_K) \cdot \mathbf{x} + (w_{i0} - w_{K0}))}$$

■ Substitute in expression (5):

$$p_y(y = i|\mathbf{x}) = \frac{\exp((\mathbf{w}_i - \mathbf{w}_K) \cdot \mathbf{x} + (w_{i0} - w_{K0}))}{1 + \sum_{j=1}^{K-1} \exp((\mathbf{w}_j - \mathbf{w}_K) \cdot \mathbf{x} + (w_{j0} - w_{K0}))}$$

■ More compactly, for all i :

$$p_y(y = i|\mathbf{x}) = \frac{\exp(\mathbf{w}_i \cdot \mathbf{x} + w_{i0})}{\sum_{j=1}^K \exp(\mathbf{w}_j \cdot \mathbf{x} + w_{j0})}$$

Aside: Softmax Regression

- This is the **softmax function**
- We can use it to state that if:

$$p_y(y = k|\mathbf{x}) > p_y(y = j|\mathbf{x})$$

- Then:

$$\frac{\exp(\mathbf{w}_k \cdot \mathbf{x} + w_{k0})}{\sum_{i=1}^K \exp(\mathbf{w}_i \cdot \mathbf{x} + w_{i0})} > \frac{\exp(\mathbf{w}_j \cdot \mathbf{x} + w_{j0})}{\sum_{i=1}^K \exp(\mathbf{w}_i \cdot \mathbf{x} + w_{i0})}$$
$$\mathbf{w}_k \cdot \mathbf{x} + w_{k0} > \mathbf{w}_j \cdot \mathbf{x} + w_{j0}$$
$$f_k(\mathbf{x}) > f_j(\mathbf{x})$$

Aside: Softmax Regression

- So, the condition for the Bayes Optimal Classifier allows us to write:

$$f^*(\mathbf{x}) = k \quad \text{if:} \quad f_k(\mathbf{x}) > f_j(\mathbf{x}) \quad \forall j \neq k$$

Where $f_i(\mathbf{x}) = \mathbf{w}_i \cdot \mathbf{x} + w_{i0}$

- And this is the K Class Linear Discriminant definition

Aside: The Multiclass SVM

- Alternatively, Weston & Watkins ('99) proposed the **multiclass SVM**
- This is a formulation of the SVM that enables a multiple classification problem to be solved in a single optimisation
- This approach also gives rise to a K Class Linear Discriminant

Lecture Overview

- 1 Lecture Overview
- 2 Support Vector Classification
- 3 Multi-Class Support Vector Classification
- 4 Support Vector Regression**
- 5 Summary

Support Vector Regression

- The **SVR** follows similar principles to the SVC
- It is a learning algorithm that emerges from the optimisation of a **generalisation bound** on the **ϵ -insensitive loss**
- This is a **convex, sparsity inducing**, regression loss function

Support Vector Regression

- The optimisation problem for the linear SVR is:

$$\min_{\mathbf{w}, b, \zeta, \tilde{\zeta}} \quad \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n (\zeta^{(i)} + \tilde{\zeta}^{(i)})$$


$$\begin{aligned} \text{subject to:} \quad & y^{(i)} - \mathbf{w} \cdot \mathbf{x}^{(i)} - b \leq \epsilon + \zeta^{(i)} \\ & \mathbf{w} \cdot \mathbf{x}^{(i)} + b - y^{(i)} \leq \epsilon + \tilde{\zeta}^{(i)} \\ & \zeta^{(i)}, \tilde{\zeta}^{(i)} \geq 0 \end{aligned}$$

Here ϵ , C are **hyperparameters**

C is the trade-off parameter which controls the width of **margin** versus the **tolerance for error**

- What is the nature of the errors?


Support Vector Regression



- $L^{(i)} = \max(|y^{(i)} - \mathbf{w} \cdot \mathbf{x}^{(i)} - b| + \epsilon, 0)$

- $\zeta^{(i)}, \tilde{\zeta}^{(i)}$ are **slack variables**

Support Vector Regression



■ Here, for $\epsilon = 0$, as $C \rightarrow \infty$, we tend to **Laplace regression**

Lecture Overview

- 1 Lecture Overview
- 2 Support Vector Classification
- 3 Multi-Class Support Vector Classification
- 4 Support Vector Regression
- 5 Summary**

Summary

- 1 Adding **kernels** to the **linear SVM** allows us to build highly effective **non-linear SVM** classifiers
- 2 It is possible to extend the **binary** approach to classification to **multi-class problems** via heuristics such as **OvO** and **OvA**...with some problems
- 3 The SVM framework can be extended to tackle **regression** problems via the **SVR**. This algorithm shares the **sparsity** inducing properties of the original SVM.